

MAKE IT LOOK GOOD

Generate Professional Slides, Documents & Images With
AI

The non-designer's playbook for clean decks, polished docs, and the right
image, plus the one limitation nobody warns you about.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

Here is the sentence I hear most from smart professionals: "I know what I want it to say, I just can't make it look good, and it takes me forever."

You are not bad at your job. You were simply never trained as a designer, and design tools were built for people who were. So you lose an afternoon nudging text boxes in a slide deck, fighting with margins in a document, and giving up on the image entirely because the stock photo options all look like a 2011 brochure. The result is work that is smart on the inside and amateur on the outside, and you can feel the gap.

The good news as of 2026 is that the visual layer is the part AI has gotten genuinely good at. The same models you already use for writing can now draft a structured slide deck, format a messy document into something clean, and generate a custom image in seconds. The catch is that nobody hands you the operating manual, and there is one specific thing these tools are still bad at that will embarrass you if you do not know about it.

This guide is the manual. It covers how to produce professional slides, documents, and images with AI, where each tool genuinely shines, and the one honest limitation that separates people who get clean output from people who get garbage. Everything here reflects what these tools can actually do right now, not what a demo video promised two years ago. No design background required, and not a line of code in sight.

1. Build the deck in two passes, not one

The mistake almost everyone makes with AI slides is asking for the finished presentation in a single shot. You type "make me a 15-slide deck on our Q3 results," you get fifteen slides of bland filler, and you conclude AI can't really do this. The problem was the order of operations, not the tool.

Professional decks are built in two passes, and AI is excellent at both as long as you keep them separate. The first pass is the outline, which is the thinking. The second pass is the design, which is the polish. When you mash them together, you get a confident-looking deck full of shallow content, because the tool spent its effort making boxes pretty instead of making the argument land.

Pass one, the outline. Before you touch any slide tool, open a regular chat assistant and have it build the skeleton. Tell it the audience, the goal of the meeting, and the single decision you want people to walk out having made. Ask for a slide-by-slide outline with a

one-line headline for each slide and three supporting bullets underneath. Read it. Fix the logic. This is where you earn the deck, and it takes ten minutes.

Pass two, the design. Now you hand that approved outline to a tool that actually builds slides. In 2026 the strong options are dedicated deck generators like Gamma, which build a whole formatted deck from your outline in about two minutes, and the slide tools now baked into the big platforms, where Microsoft's Copilot builds inside PowerPoint and Google's Gemini builds editable Google Slides decks from a single prompt or document. Paste your outline, pick a clean theme, and let it lay everything out.

The honest truth about pass two is that these tools impose their own design taste. You are refining someone else's layout rather than building your own from scratch. That is a feature, not a bug, when you are not a designer, because their default taste is better than what most of us would assemble by hand. You just need to know going in that your job is to edit and trim, not to accept the first draft as final.

Before: A blank presentation file, a 9 a.m. board meeting, and the familiar dread of building 20 slides from nothing. Two hours of dragging text boxes ahead of you, and the design will still look like a template.

After: Ten minutes spent perfecting a tight outline in a chat assistant. The headline of every slide already says something, because you wrote the argument first. You paste it into a deck builder, choose a restrained theme, and ninety seconds later you have a formatted draft. You spend the next twenty minutes cutting weak slides and fixing two charts. Done in well under an hour, and it argues a clear point instead of just listing facts.

The deck looks good because the structure was good before the design ever touched it. That sequence is the whole trick.

2. Refine the deck like a picky editor, not a beggar

Once you have a generated deck, the difference between amateur and professional output is entirely in the refinement, and refinement is a skill you already have from editing your own writing.

Most people, faced with an AI-built deck, either accept it as-is or give up and rebuild it manually. Both are wrong. The move is to talk to the deck the way a sharp manager talks to a junior who did a decent first draft. You are specific, you are demanding, and you point at exactly what is off.

The newer deck tools let you request changes in plain language across the whole presentation at once. Instead of fixing one slide at a time, you say what you want and it propagates. So your refinement pass sounds like a series of direct notes.

"Slide 4 is too crowded. Split it into two slides, one for the problem and one for the cost of the problem."

"Every slide has five bullets. Cut each one down to a maximum of three, and make the first word of each bullet a verb."

"The tone is too cheerful for a budget meeting. Make the language plainer and more serious."

"Turn the data on slide 8 into a simple bar chart instead of a bulleted list."

Each of these is a note you could give a person. The reason this works better than rebuilding is that you are spending your judgment, which is your real value, instead of your time pushing pixels. Three or four rounds of pointed notes will take a generic generated deck and make it specifically yours.

One caveat worth knowing. These tools are far better at refining content that already exists than at inventing missing content from a cold start. If a slide is empty because you never gave it substance, no amount of "make this better" will fill it. Go back to your outline and give it the substance first.

3. Turn messy notes into a clean document

The same two-pass logic rescues you from the other visual nightmare, which is the document that is full of good thinking and looks like it was typed in a hurry on a phone. Meeting notes, a rough proposal, a policy draft, a long email you want to turn into a one-pager. The content is there. The formatting is a disaster.

This is one of the most reliable wins available, because formatting a document is a structured, rule-based job and AI is genuinely strong at structured, rule-based jobs. You are not asking it to be creative. You are asking it to impose order on something you already wrote.

The workflow is simple. Paste your raw notes and tell the assistant exactly what finished document you want them to become. Be specific about the structure, because structure is what makes a document read as professional. Ask for clear section headings, short paragraphs instead of walls of text, a summary at the top, and a consistent format for any

lists. If it is going to a specific audience, say so, because a document for your team and a document for a client get organized differently.

Before: Forty lines of bullet fragments from a strategy session. Half are complete thoughts, half are three words you scribbled to remember a point. You know there is a solid proposal buried in there, but turning it into something you would actually send your director feels like an hour of reformatting you keep postponing.

After: You paste the whole mess into the assistant with one instruction: "Turn these raw notes into a one-page proposal for my director. Open with a two-sentence summary of the recommendation. Use three clear sections with headings: the problem, the proposed solution, and what we need to approve it. Keep paragraphs to three sentences. Plain, confident tone." Fifteen seconds later you have a clean draft. You spend five minutes correcting two facts it smoothed over, and it is ready.

A word of honest caution that matters more for documents than for slides. The assistant will sometimes smooth a rough note into a confident sentence that says slightly more than you meant, or fills a gap with a reasonable-sounding detail you never provided. Always read the formatted version against your original notes and check that it did not invent specifics. The formatting is trustworthy. The facts are your job to verify.

4. Generate an image that keeps your real subject

Images are where most people either give up or get burned. The giving-up version is settling for a generic stock photo. The getting-burned version is generating something that looks impressive but is subtly wrong, like a product that resembles yours but is not actually yours.

Modern image tools split cleanly into two jobs, and knowing which job you are doing changes how you prompt.

The first job is a generic illustration where the exact subject does not matter. A picture of "a busy modern office," "an abstract background in blue and slate," "a friendly illustration of two people shaking hands." For these you just describe what you want in plain, specific language. Name the setting, the mood, the color palette, and the style, the same way you would brief a photographer. The more concrete you are, the closer the first result lands.

The second job is the one people get wrong. It is when the image must contain a specific real thing. Your actual product. A real person. Your physical storefront. If you only describe it in words, the tool will invent a plausible lookalike, and a plausible lookalike of your own

product is worse than useless. The fix, as of 2026, is to attach a reference image. The current generation of tools, like Google's Nano Banana models, are built specifically to take a reference photo and preserve that subject's real likeness across new scenes and edits. You upload a clear photo of the actual thing, then describe the new setting you want it placed in.

Before: You need a header image of your product on a clean desk for a landing page. You describe it in detail to an image tool. What comes back looks slick and is completely wrong, a generic gadget that shares your product's category but none of its actual features. Sending a customer a picture of a product you do not sell is not an option, so you fall back to the same tired stock photo as everyone else.

After: You take one clear, well-lit photo of your actual product with your phone. You upload it to the image tool and write: "Place this exact product on a clean light-wood desk by a window, soft morning light, a notebook and a coffee cup slightly out of focus in the background, photographed from a low angle." The result keeps your real product recognizable and drops it into a setting that looks professionally shot. It is your product, in a scene you could never afford to stage.

The rule is short and worth memorizing. If it must be your real subject, attach a reference photo. If any version of the subject will do, describe it in words. Most of the frustration people have with AI images comes from trying to describe a specific real thing in words and getting a confident stranger back.

5. The one honest limitation: AI cannot spell inside an image

This is the single most important thing in this guide, and almost nobody tells you about it before you waste an hour.

AI image tools are bad at putting readable text inside an image. You ask for a poster with your event name on it, or a logo with your company name, or a social graphic with a headline, and the words come back misspelled, melted, or turned into letter-shaped gibberish. This is not you using it wrong. It is a real, known limitation of how these tools work.

Here is the plain reason. Image models generate what text looks like, not what text says. They were trained to match the overall look of a picture, so to the model your headline is just another visual texture sitting in the scene, like a brick pattern or the grain of wood. It

paints something that statistically resembles letters in that spot. It has no idea that those letters have to spell a specific word. The tools have improved at this in 2025 and 2026, and a short clean word on a simple background sometimes comes out right. But for anything text-heavy or anything where a misspelling would embarrass you, you cannot trust it.

The fix is simple once you accept the division of labor. Use AI for the picture, and use a real design tool for the words. Generate your background, your photo, or your illustration with the image AI, then drop that image into a layout tool like Canva, PowerPoint, or Google Slides and type the text on top yourself with a normal text box. The text tool spells correctly every time because it is using actual fonts, not painting pixels. You get the best of both, a custom AI visual underneath and crisp, correct, on-brand words on top.

Before: You generate a beautiful event poster with the title baked into the image by the AI. The art is genuinely good. The title reads "ANUUAL LEADRSHIP SUMMMIT." You try ten more times with the spelling spelled out in the prompt, in all caps, in quotes. It garbles a different word each time. An hour gone, still no usable poster.

After: You generate the poster as a clean background only, no text at all in the prompt. You get a striking image in one try. You open Canva, drop the image in as the background, and add the real title, date, and venue in proper text boxes with a font you like. Five minutes, perfectly spelled, fully editable later when the venue changes.

Memorize this division and you will never lose another hour to garbled text. AI makes the picture. A layout tool puts the words on. They are good at different halves of the same job.

6. The repeatable workflow

Here is the whole thing as one loop you can run every time, whether the output is a deck, a document, or an image. It is the same shape each time, which is what makes it fast once it is a habit.

Step one, decide the job and the audience. One sentence. "A deck to get the board to approve the budget." "A one-page proposal for my director." "A header image of my real product for the landing page." Naming the audience and the goal up front is what keeps the output from being generic.

Step two, do the thinking before the polish. For a deck, write the outline first. For a document, gather your raw notes. For an image, decide whether the subject must be your

real thing, which tells you whether to attach a reference photo. The thinking is your job and AI cannot do it for you.

Step three, generate the first draft. Hand the thinking to the right tool. The outline goes to a deck builder. The notes go to an assistant with clear formatting instructions. The image description, plus a reference photo if the subject is real, goes to an image tool. Expect a rough but useful first version.

Step four, refine with specific notes. Talk to the draft like a picky editor. Point at exactly what is wrong and say exactly what you want instead. Three or four pointed rounds beats one vague "make it better."

Step five, add real text in a real tool. If the output is an image that needs words on it, that text goes on in a layout tool, never inside the image generator. This is the step that keeps you from shipping something embarrassing.

Step six, fact-check before you send. Read the polished version against your original substance. Confirm it did not invent a detail or smooth a rough fact into a confident wrong one. The look is trustworthy. The facts are yours to verify.

Run this loop a few times and it stops feeling like a process and starts feeling like how you work. The afternoon you used to lose to formatting becomes forty minutes, and the output looks like someone with a design team made it.

The one-page cheat sheet

Keep this beside you. The order matters as much as the steps.

1. **Two passes for decks.** Outline first in a chat assistant, design second in a deck builder. Never both at once.
2. **Refine, don't rebuild.** Give the generated deck specific notes like a picky editor. Spend judgment, not pixel-pushing time.
3. **Documents are a formatting job.** Paste raw notes, ask for headings, short paragraphs, a top summary, and a named audience.
4. **Images split in two.** Generic subject, describe it in words. Your real subject, attach a reference photo so it keeps the real likeness.
5. **AI cannot spell in images.** Generate the picture with no text, then add the real words in Canva, PowerPoint, or Slides with a normal text box.
6. **The loop:** job and audience, then thinking, then first draft, then specific-note refinement, then real text in a real tool, then fact-check.

7. Trust the look, verify the facts. Formatting and layout are reliable. Invented details are not. Always read the polished version against your source.

You were never bad at this. You were missing the manual and the one limitation nobody mentions. Now you have both.

*This is a companion to **The Artificial Advantage: Mastering AI in the Modern Workplace** by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%. Get it at elvinpeters.com.*

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google, OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.