

DECODED

The 8 Pro AI Techniques Hidden in the Developer Docs

The prompt-engineering best practices Anthropic, Google, and OpenAI publish for their engineers, in plain English for professionals.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

The people getting astonishing results from AI are not smarter than you. They have read something you have not.

Every major AI company publishes a set of "prompt engineering" best practices. These guides are written for software developers, buried in documentation sites, and wrapped in jargon. Almost no professional outside of engineering ever reads them. That is the gap. The exact same techniques that engineers use to build reliable AI products work just as well when you are writing an email, building a report, or analyzing a spreadsheet.

This guide takes the eight techniques Anthropic documents (and that Google and OpenAI echo almost word for word) and strips out the code. What is left is eight concrete moves you can use today, in any tool, with no technical background. They are listed in the order the pros apply them, from the simplest to the most advanced.

A note on truth: there is a lot of garbage AI advice online. Everything here is drawn from the official developer documentation of the companies that actually build these models, not from someone's hot take. These techniques are stable. They worked last year, they work now, and they transfer across ChatGPT, Claude, Gemini, and Copilot.

1. Be clear and direct

What the docs call it: "Be clear and direct."

What it actually means: Tell the AI exactly what you want, the way you would brief a sharp new hire on their first day. That person is capable, but they cannot read your mind and they have no idea what you already know.

Most weak results come from prompts that assume context the AI does not have. You know the client is sensitive about deadlines. You know the report goes to the board. The AI knows none of that until you say it.

How to do it: Spell out the task, the audience, the format, the length, and the tone. If a detail would change the answer, include it.

Vague: "Write an email about the project delay."

Clear and direct: "Write a short email to a client, Sarah Chen, telling her our project will finish two weeks late, on March 15 instead of March 1. She values transparency and

hates being managed. Be straightforward, take ownership, and keep it under 150 words. Do not over-apologize."

The second prompt does not require more skill. It requires one extra minute of saying what you already know.

2. Show it examples

What the docs call it: "Multishot prompting" or "using examples."

What it actually means: Showing beats telling. If you give the AI two or three examples of what a good result looks like, it will match the pattern far more reliably than if you describe the pattern in words.

This is the single highest-leverage technique most people skip. Models are pattern-matchers at their core. Hand them the pattern.

How to do it: Paste a few examples of the output you want before you give it the real task. They can be your own past work.

"Here are three subject lines that performed well for us: 1. The one mistake costing you an hour a day 2. Your competitor just did this (and it worked) 3. We tried it so you don't have to

Now write five more subject lines in that same style for a webinar on AI at work."

The AI now knows your voice, your length, and your angle, because you showed it instead of describing it.

3. Let it think first

What the docs call it: "Chain of thought," or "letting Claude think."

What it actually means: When a task involves reasoning, math, comparison, or judgment, tell the AI to work through it step by step before giving the final answer. Quality goes up measurably when it reasons out loud first.

The reason is simple. A model that jumps straight to an answer is guessing. A model that lays out its steps catches its own mistakes along the way, the same way you would if you

showed your work.

How to do it: Add a line like "Think through this step by step before you answer," or "First list the factors, then weigh them, then give your recommendation."

"I have three vendors with different pricing and contract terms (pasted below). Before recommending one, walk through the trade-offs of each step by step. Then give me your recommendation and the single biggest risk."

For anything where being wrong has a cost, this one technique is worth the extra few seconds it takes the AI to respond.

4. Label your sections

What the docs call it: "Using XML tags."

What it actually means: When your prompt has different parts (instructions, background, the data to work on, examples), label each part clearly so the AI never confuses your instructions with your content. Engineers use little tags like `<instructions>` and `<data>`. You do not need the tags. You need the habit of labeling.

This matters most when you paste in a long document and then ask a question about it. Without labels, the AI can mistake your question for part of the document, or answer the document instead of you.

How to do it: Put a clear header before each block. Plain words work perfectly.

INSTRUCTIONS: Summarize the feedback below into the three most common complaints. Quote one real example for each.

CUSTOMER FEEDBACK: [paste the 40 reviews here]

The AI now knows exactly what is the job and what is the material. This small habit removes a whole category of "it answered the wrong thing" frustration.

5. Give it a role

What the docs call it: "Giving Claude a role," or "system prompts."

What it actually means: Tell the AI who to be before you tell it what to do. "You are a senior financial analyst" produces a different, sharper answer than no role at all, because it pulls the model toward the right knowledge, vocabulary, and standards.

A role is the fastest way to raise the quality and the expertise of an answer in a single sentence.

How to do it: Start with one line that assigns expertise and a standard.

"You are a veteran hiring manager who has reviewed 10,000 resumes and has zero patience for fluff. Review the resume below and tell me, honestly, the three things a recruiter would hold against it in the first ten seconds."

The role does two jobs. It raises the expertise, and it sets the attitude. "Honest and impatient" gets you a more useful answer than "helpful and encouraging."

6. Start its answer for it

What the docs call it: "Prefilling the response."

What it actually means: You can put the first few words of the answer into the AI's mouth to steer the format and skip the throat-clearing. If you want a table, start the answer with the table header. If you want it to skip the "Sure, here's..." preamble, start the real answer yourself.

This is a quiet pro move that most tools support through a simple trick: just tell it how to begin, or begin for it.

How to do it: Tell it the exact first words, or the exact format to open with.

"List the pros and cons of remote work for our team. Start your response directly with a markdown table with two columns, Pros and Cons. No introduction, no summary, just the table."

You get clean, usable output instead of a paragraph of warm-up you have to delete.

7. Break big jobs into steps

What the docs call it: "Prompt chaining."

What it actually means: Do not ask the AI to do a large, multi-part job in one shot. Break it into a sequence, where each step uses the result of the last. The output of each step is cleaner, and you can catch a problem early instead of at the end.

Asking for everything at once is how you get a mediocre version of five things. Chaining gets you a strong version of each.

How to do it: Run it as a conversation, one move at a time.

Step 1: "Here are my rough notes from the client call. Pull out every commitment we made and every open question." *(check the result)*

Step 2: "Good. Now turn the commitments into a follow-up email to the client." *(check the result)*

Step 3: "Now turn the open questions into a task list for my team."

Each step is small, correct, and easy to fix. The final result is built on solid pieces instead of one big guess.

8. Feeding it long documents

What the docs call it: "Long context tips."

What it actually means: When you give the AI a long document to work from (a contract, a report, a transcript), the order matters. Put the long document near the top, and put your actual question at the very bottom. Then ask it to find and quote the relevant parts before it answers.

Models pay the most attention to the start and the end of what you give them. Putting your question last keeps it front of mind. Asking it to quote first keeps it honest, because it has to point to the actual text instead of making something up.

How to do it:

[paste the full 20-page contract first]

"Using only the contract above, answer this: what is our notice period to cancel, and are there any penalties? First quote the exact clauses you are relying on, then give me a plain-English answer."

The "quote first" instruction is the difference between a confident wrong answer and a sourced right one.

The one-page cheat sheet

Keep this beside you. Stack these moves, do not pick just one.

1. **Be clear and direct.** Brief it like a new hire. Say the audience, format, length, tone.
2. **Show examples.** Two or three samples of good beats any description.
3. **Let it think.** "Work through this step by step before answering."
4. **Label your sections.** INSTRUCTIONS here, DATA there. Never let it confuse them.
5. **Give it a role.** "You are a [expert] with [standard]." Sets expertise and attitude.
6. **Start its answer.** Tell it the first words or the exact format. Skip the preamble.
7. **Break big jobs into steps.** A strong version of each beats a weak version of all.
8. **Long documents:** document first, question last, "quote the relevant parts first."

These eight are not tricks. They are how the tool was built to be used. Most people will never read the manual. You just did.

This is a companion to [The Artificial Advantage: Mastering AI in the Modern Workplace](#) by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%. Get it at elvinpeters.com.