



FEED IT YOUR WORLD

The Context Playbook: How to Make AI Work With Your
Documents, Data & Details

The reason AI answers feel generic is simple. You never gave it your world
to work from. Here is how to fix that.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

Most people who say "AI gives me generic advice" are right, and the reason is not the model. It is that they keep asking a brilliant stranger to comment on a business it has never seen. The AI does not know your client, your numbers, your brand voice, your last six months of email threads, or the contract sitting in your downloads folder. So it answers the average version of your question, because the average is all you gave it.

The fix is not a better prompt. It is better context. The single biggest difference between someone who gets bland output and someone who gets answers that sound like they came from inside their own company is that the second person feeds the AI their actual material first.

This guide is about that one skill. How to give AI your documents, your data, and your details so its answers are about your situation instead of a generic one. There is a craft to it. Where you put the document in the prompt changes the answer. How much you paste changes the answer. Whether you set the context up once and reuse it, or re-explain yourself every single time, changes your whole week.

A note on truth, because there is a lot of confident nonsense online about this. Everything here is grounded in the official documentation from the companies that build these tools, including the prompt-engineering guidance Anthropic publishes for developers, the help docs for custom assistants from OpenAI and Google, and the published limits for Google's research tool NotebookLM. Where I cite a number or a behavior, it comes from the source, not from a hunch.

1. Stop asking. Start showing.

The mental shift that fixes most generic answers is this. The AI is not a search engine you query. It is a sharp analyst who knows nothing about your specific situation until you hand it the file.

When you ask "how should I price my consulting package," you get a textbook answer, because a textbook is all the AI has to draw on. When you paste your current rate card, your three competitors' pricing, and a note about the client you are quoting, then ask the same question, you get a recommendation built on your reality.

The pattern is always the same. Give it the raw material first, then ask the question. Most people do the opposite. They ask the question and expect the AI to somehow already know the material.

How to do it: Before you type your question, ask yourself one thing. What would a smart human need to see to answer this well? A contract. Last quarter's numbers. The email thread. Your style guide. Whatever it is, paste it in first. The answer quality tracks almost perfectly with the quality of what you feed it.

Generic ask: "Write a follow-up email to a prospect who went quiet."

Fed its world: "Below is my last email to this prospect and her reply from three weeks ago. She runs operations at a mid-size logistics firm and mentioned budget approval was stuck with her CFO. [paste the thread] Now write a short, warm follow-up that gives her an easy reason to re-engage without sounding desperate."

The second email lands because it knows the CFO is the blocker. You knew that. Now the AI does too.

2. How much context is enough

The instinct most people have is to give the AI a one-line summary because they think a wall of text will confuse it. That instinct is backwards. These tools are built to read large amounts of material, and they get sharper, not dumber, when you give them the real thing instead of your compressed version of it.

There is a real limit, called the context window, which is the total amount of text the AI can hold in mind at once. For everyday work in the major chat tools, that window is large. You can comfortably paste a long contract, a full meeting transcript, or a dozen customer reviews and still have room. You will know you have hit the ceiling only on genuinely huge inputs, like an entire book or hundreds of pages at once, and the tool will usually tell you.

So the practical rule is the opposite of what people fear. Give it more, not less. Paste the whole email thread, not your summary of it. Include the boring parts. The AI will filter what matters, and the details you would have trimmed are often the ones that change the answer.

How to do it: When in doubt, include it. The cost of pasting too much is small. The cost of pasting too little is a generic answer. The only thing to avoid is pasting many unrelated topics into one chat, which muddies the focus. Keep one chat to one job.

Too little: "I had a tough performance review conversation with a team member. How do I follow up?"

Right amount: "Here are my actual notes from the review meeting, including the three concerns I raised and how she responded to each. [paste the full notes] She got defensive on the second point about deadlines. Draft a follow-up message that acknowledges her perspective, restates the one non-negotiable, and keeps the door open."

The full notes carry the emotional texture and the specifics. That is exactly what the AI needs to write something that sounds like you were actually in the room.

3. The right way to feed a long document

This is the technique that separates people who get reliable answers from documents and people who get confident, wrong ones. When you hand the AI a long file, a contract, a report, a transcript, the order you put things in genuinely changes the result.

Anthropic states this directly in its published prompt-engineering guidance. Put the long document near the top of your prompt, above your question and your instructions. Their words are to "place your long documents and inputs near the top of your prompt, above your query." The reason is that the models pay extra attention to material placed early, and putting your actual question last keeps it fresh in mind when the answer gets written.

There is a second move that matters even more. Ask the AI to quote the relevant parts of the document before it answers. Anthropic's guidance again, almost verbatim, is to "ask Claude to quote relevant parts of the documents first before carrying out its task." This forces the AI to point at real text instead of inventing a plausible-sounding answer. A model that has to quote the actual clause cannot easily make one up.

How to do it: Three steps, every time you work from a long document.

1. Paste the full document first, at the top.
2. Put your question at the very bottom.
3. Tell it to quote the relevant passages before answering.

[paste the full 18-page vendor contract first]

"Using only the contract above, answer two questions. What is our notice period to terminate, and are there any auto-renewal penalties? First quote the exact clauses you are relying on, word for word. Then give me a plain-English answer for each. If the contract does not address something, say so rather than guessing."

The "quote first" line and the "say so rather than guessing" line together are the difference between a sourced answer you can trust and a smooth answer you cannot. This works in every major chat tool, whether you paste the text or attach the file.

4. Set up your context once, reuse it forever

Here is the move that gives back hours. If you find yourself re-explaining the same background at the start of every conversation, your company, your role, your brand voice, your standard rules, you are doing manual labor the tool was built to eliminate.

Every major chat tool now lets you save persistent context that loads automatically into a dedicated workspace, so you never paste it again. The names differ, the idea is identical.

In the tool from Anthropic, these are called Projects. Per its support documentation, a Project is a "self-contained workspace" with its own knowledge base and custom instructions, where you "upload relevant documents, text, code, or other files" that the AI uses for context across every chat in that workspace. On paid plans, the documentation notes the knowledge base scales to hold large amounts of content.

In the tool from OpenAI, these are custom GPTs. Per its help center, you give a GPT standing instructions for how it should behave, and you upload knowledge files, which are "reference material you want the GPT to draw from when answering." You can attach up to twenty such files. The docs add a useful tip: if you want it to quote or cite the uploaded material, say so in the instructions.

In the tool from Google, these are called Gems. Per Google's guidance, a Gem bundles a name, a set of instructions, and up to ten knowledge files into a reusable assistant, and Google recommends writing instructions of roughly 500 to 2,000 characters. A notable perk for Google Workspace users is that a Gem can connect directly to Drive files that update automatically, so your assistant always reads the current version.

How to do it: Pick the tool you already use. Create one workspace for a recurring job. Load it with the documents and the standing instructions that job always needs. Then every chat inside it starts already knowing your world.

What you stop typing every time: "I run marketing for a B2B software company. Our voice is direct and a little irreverent, never corporate. We sell to operations managers, not executives. Avoid buzzwords like synergy and leverage..."

What you set up once instead: A workspace named "Marketing Copy" with a brand voice document and a one-paragraph instruction saved as its standing context. Now you

just open it and say "draft three LinkedIn posts about our new reporting feature," and it already knows the voice, the audience, and the rules.

That is the same prompt producing a far better result, because the context lives in the room instead of being re-explained at the door.

5. Build reusable context blocks you can paste anywhere

Persistent workspaces are powerful, but you will not always be inside one. Sometimes you are in a quick chat, a different tool, or a colleague's screen. For those moments, keep a small library of reusable context blocks in a plain notes file or document.

A context block is a short, pre-written paragraph that establishes your world, ready to paste at the top of any conversation. Think of it as a backstage briefing you hand the AI before the real request. Writing it once and reusing it beats re-typing it badly under time pressure.

The blocks worth keeping are the ones you would otherwise explain over and over. Your company and role in three sentences. Your brand voice with two dos and two don'ts. Your ideal customer. Your standard output preferences, like "always give me a draft I can edit, never a finished thing you assume is right."

How to do it: Open a notes file. Write three or four of these blocks today. Give each a clear label. When you start a fresh chat anywhere, paste the relevant block first, then make your request.

A saved "Who We Are" block:

"Context for this conversation: I run a boutique accounting firm serving small business owners in the trades, plumbers, electricians, contractors. My clients are smart but not financially fluent and they are allergic to jargon. My voice is plain, warm, and reassuring. When you write for them, explain every term in plain English and never assume they know an acronym."

Then your actual request: "Using that context, write a short email explaining what a quarterly tax instalment is and why missing one is a problem."

The same email request, with and without that block in front of it, produces two completely different results. One sounds like a tax textbook. The other sounds like you talking to your clients.

6. Use a tool built to live inside your own files

For the specific job of working deeply with a fixed set of your own documents, there is a category of tool built for exactly that, and the best-known example is Google's NotebookLM. It is worth knowing about even if you live in the chat tools, because it solves a different problem.

A normal chat tool draws on everything it learned in training, plus whatever you paste. NotebookLM does the opposite. It answers only from the sources you upload, and nothing else. Google's documentation describes it as a closed system grounded in your material. You can load up to fifty sources into a single notebook, including PDFs, documents, websites, and transcripts, with each source holding up to 500,000 words.

The feature that makes it trustworthy is the citations. Every answer points back to the exact passage in your documents that it came from, with the page or location, so you can verify it in one click. This is the cure for the deepest version of the "I cannot trust the answer" problem. The AI is not allowed to wander off into general knowledge. It has to stay inside the binder you gave it.

How to do it: Use this kind of tool when you have a defined body of material you will question repeatedly. A research folder, a stack of policy documents, a long report, the transcripts from a dozen customer interviews. Load the sources once, then ask it questions and follow the citations to confirm.

The situation: You have eleven recorded customer interviews, transcribed, and you need to find the patterns without re-reading all of them.

The move: Upload all eleven transcripts as sources. Then ask: "Across these interviews, what are the three most common frustrations customers mentioned about our onboarding? For each one, cite the specific interviews and quote the lines where it came up."

You get a themed summary where every claim links back to a real customer saying a real thing. That is an analysis you can put in front of your boss, because you can prove every line of it.

7. Keep your context current, and know what not to feed it

Context is only useful if it is true. The fastest way to make a persistent workspace turn against you is to load it with last year's pricing, an old brand voice, or a policy that has since changed, then forget you did. The AI will faithfully use the stale material and give you confidently outdated answers.

Treat your saved context like a pantry. Things expire. When your pricing changes, update the document in the workspace. When your positioning shifts, refresh the standing instructions. Google's Gems even let Workspace users connect live Drive files that update on their own, which sidesteps the staleness problem for the documents you change most.

There is also a line you should not cross. Be deliberate about what you paste into these tools, especially material that is confidential, regulated, or covers other people's private information. Client financial records, employee health details, anything under a strict confidentiality agreement, these deserve real thought about your company's policy and the tool's data terms before they go anywhere near a chat box. The point of this guide is to feed the AI your world. Feeding it responsibly is part of the craft.

How to do it: Once a month, glance at your saved workspaces and context blocks and ask "is any of this out of date." Update what changed. And before pasting anything sensitive, ask one question. Would I be comfortable if this material left the building. If the answer is no, redact it or keep it out.

Stale and dangerous: A "Sales" workspace still loaded with a discount policy that ended in March, plus a customer list with full credit card numbers pasted in for "context."

Current and clean: The same workspace with the live pricing document connected so it auto-updates, the expired discount removed, and the customer examples stripped of any payment or personal data before they went in.

Good context is a living thing. Tend it, and keep the genuinely sensitive material on the right side of the line.

The one-page cheat sheet

Keep this beside you. The whole game is feeding the AI your world before you ask it anything.

1. **Show, don't ask.** Paste the real material first, then your question. Generic in, generic out.
2. **Give it more, not less.** These tools handle long inputs. Paste the whole thing, not your summary. One chat, one job.
3. **Long documents:** put the document at the top, your question at the bottom, and tell it to quote the relevant parts before it answers.
4. **Set context once.** Use a persistent workspace (Projects, custom GPTs, or Gems) loaded with your standing documents and rules. Stop re-explaining yourself.
5. **Keep paste-ready blocks.** A notes file with three or four pre-written "who I am" paragraphs you drop into any fresh chat.
6. **For a fixed set of files,** use a tool built to live inside them, like NotebookLM, where every answer cites the exact source.
7. **Keep it current and keep it safe.** Update stale context monthly. Think before pasting confidential or personal data anywhere.

The difference between generic answers and answers that sound like they came from inside your own company is not a secret prompt. It is context. Most people never give it. You now know how.

This is a companion to [The Artificial Advantage: Mastering AI in the Modern Workplace](#) by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%. Get it at elvinpeters.com.

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google, OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.