STAY AHEAD

Your 30-Minutes-a-Week AI Operating System

The simple weekly routine that keeps your AI skills current as the models change, so your edge compounds instead of going stale.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

The most common fear professionals have about AI is not "I can't learn it." It is "I'll learn it, and then it'll change, and I'll be behind again." That fear is reasonable. The tools really do move fast. A model you mastered in spring behaves differently by autumn, a feature you relied on gets replaced, and the prompt that worked beautifully last quarter now produces something slightly worse or slightly better and you cannot tell which.

Here is the part nobody tells you. The people who stay ahead are not the people who chase every announcement. They are the people who have a small, boring, repeatable routine. Thirty minutes a week, the same handful of moves, done consistently. That is enough to keep your edge sharp while everyone around you swings between panic and neglect.

This guide gives you that routine. It is not the practice plan from *The Artificial Advantage*, which is about building your foundational skill from zero. This is the maintenance layer that runs forever after. Think of it as an operating system for your AI ability. You install it once, it runs quietly in the background, and it keeps everything else you have learned from rotting.

A note on truth before we start. The biggest danger in staying current is not missing news. It is drowning in it. Most of what gets posted about AI is hype, recycled press releases, or someone trying to sell you a course. The routine below is built to filter that noise hard, so the thirty minutes you spend is signal and almost nothing else.

Total time: about thirty minutes a week, plus one slightly longer session once a month. You can do the weekly block over a coffee. Pick a fixed slot, protect it, and let it run.

---

1. The weekly test: try one new thing on one real task

The move: Once a week, take one real task from your actual work and do it with AI in a way you have not tried before. Not a practice exercise. A real thing you needed to do anyway.

This is the heart of the whole system, and it is the step people skip because it feels too simple to matter. It matters more than anything else. You do not get better at AI by reading about AI. You get better by using it on work that has real stakes, where you can feel the difference between a good result and a bad one.

The "one new thing" can be small. A feature you have not opened. A different model in the dropdown. A way of structuring your request you read about. A file type you have never uploaded. The point is to stretch one inch past what you already do on autopilot, every single week. Inches add up fast.

How to do it:

1. Pick a task you genuinely have to do this week. A report, an email thread, a plan, an analysis.
2. Choose one thing to do differently. Just one.
3. Do the task that way. Keep your old method in your back pocket in case the new one flops.
4. Notice honestly whether the result was better, worse, or the same.

A professional example: Dana runs operations for a mid-size logistics firm. Every Monday she writes a status update to the leadership team, and she has always just typed it from memory. This week, her "one new thing" is to first paste last week's update plus this week's raw numbers into the AI and ask it to flag what changed and what leadership will most want to know, before she writes a word. It takes her four extra minutes. The draft it helps her produce catches a cost spike she would have buried in paragraph three. Next week she will keep that habit and try something else on top of it.

The compounding effect is the whole point. Fifty-two small experiments a year, each one keeping what worked and dropping what did not, and you end up fluent in a way no single course could ever deliver.

---

2. Keep a personal "what works for me" file

The move: Keep one simple document where you save the prompts, phrasings, and small workflows that have actually worked for you. When you hit something good, paste it in. When you need it again, copy it out.

Every professional who is genuinely strong with AI has some version of this, even if they never gave it a name. It is the single highest-return habit in this guide, because it turns every lucky win into a permanent asset. Without it, you re-solve the same problem every month and slowly forget your own best moves. With it, your skill accumulates in writing.

This is not a fancy system. A single note, a doc, a page in whatever you already use. The format does not matter. The habit of capturing does.

How to do it:

1. Open one document. Call it whatever you like. "My AI playbook" works fine.

2. Whenever a prompt or approach gives you a noticeably good result, paste it in with a one-line note on what it was for.
3. Organize it loosely by job: emails, reports, meeting prep, analysis, whatever your real categories are.
4. When you start a task, glance at the file first. Reuse beats reinventing.

A professional example: Marcus, a financial controller, kept getting great variance explanations out of the AI but could never remember exactly how he had asked. So he started a file. Now it holds his proven setup for month-end commentary, his favorite way to turn a messy spreadsheet into a board summary, and the exact role-and-tone line that makes the AI write in his voice instead of generic corporate mush. When a new junior analyst joined, Marcus handed her the file, and she was producing usable work in a day instead of a month. The file became a training tool he never set out to build.

A word on why this beats relying on the tool's memory features. Built-in memory and saved prompts inside any given app are useful, and you should use them. But they live inside that one app. When the model changes, when you switch tools, or when your company swaps vendors, your own file comes with you. It is the one thing in this whole landscape that you fully control.

---

3. Follow a small set of trustworthy sources, on purpose

The move: Pick a tiny, deliberate set of high-signal places to follow, and ignore everything else. Two or three sources, total. The goal is not to know everything. The goal is to never be blindsided.

The instinct when you feel behind is to follow more accounts, join more groups, subscribe to more newsletters. That makes the problem worse. More inputs means more noise, more hype, more anxiety, and paradoxically less real understanding. The pros do the opposite. They follow a handful of sources they trust and tune out the firehose.

There are two kinds of source worth your time, and almost nothing else is.

Kind one: the official vendor changelogs. This is the primary source, straight from the companies that build the models. No spin, no hot takes, no engagement-bait. When a feature ships or a model changes, this is where it is announced first and described accurately. Skim the one for whatever tool you actually use. These are all real, current, and free as of 2026:

- **Claude release notes** (Anthropic): the Claude Help Center release notes page, and for the apps, the support site's "what's new."
- **ChatGPT release notes** (OpenAI): the OpenAI Help Center "ChatGPT — Release Notes" page, plus the separate "Model Release Notes" page when you want to know what actually changed in the model itself.
- **Gemini release notes** (Google): the Gemini Apps release updates page at gemini.google, and the Google Workspace Updates blog if you live in Docs, Gmail, and Sheets.
- **Copilot / Microsoft 365** (Microsoft): the Microsoft 365 roadmap and the Copilot release notes in the Microsoft support and tech community pages.

You do not read these word for word. You skim the last week or two, looking only for things that touch the tools you personally use.

Kind two: one or two high-signal human filters. A good writer who reads everything so you do not have to, and tells you what actually matters for work rather than what is loud. The strongest single recommendation for a non-technical professional is **Ethan Mollick's "One Useful Thing,"** a free weekly newsletter focused squarely on what AI means for real work, written by a Wharton professor who tests things himself rather than repeating press releases. If you want one daily skim on top of that, pick exactly one aggregator such as **TLDR AI** or **The Rundown AI**, both free, both designed to be read in a couple of minutes. Pick one. They cover the same launches with different framing, so subscribing to all three just triples your inbox for no extra signal.

That is the entire media diet. Vendor changelogs for truth, one human filter for meaning, optionally one daily skim for breadth. Three sources, thirty seconds to two minutes each.

How to filter hype from truth, fast. When you see a breathless claim about AI, run it through four quick questions before you believe or repeat it:

1. **Who is saying it, and what are they selling?** A vendor's own benchmark, a course creator's thread, and an independent tester are three very different things.
2. **Is there a primary source?** If the claim cannot be traced back to a changelog, a paper, or a hands-on test, treat it as a rumor.
3. **Does it match what I see when I try it?** Your own real task is the ultimate fact-checker. A model that "changes everything" should change something in your Tuesday.
4. **Is this new, or is it recycled?** A huge share of AI "news" is the same announcement re-posted by twelve accounts for clicks.

A professional example: Priya, a marketing director, used to feel a jolt of anxiety every time a colleague forwarded an "AI just changed everything" article. Now she has a rule.

She follows only the ChatGPT release notes page (her team's main tool) and "One Useful Thing." When someone forwards hype, she checks whether it shows up in either source. Nine times out of ten it does not, and she deletes the email with a clear conscience. The tenth time, it is real, and she is the first on her team to know what it actually means rather than what the headline screamed.

---

4. The monthly re-test: check that your tools still do what you think

The move: Once a month, spend a little longer reopening your "what works for me" file and re-running two or three of your most important prompts to confirm they still produce good results. Models change underneath you silently. This is how you catch it.

This is the step almost nobody does, and it is the difference between someone who looks current and someone who actually is. When a vendor updates a model, your trusted prompts do not announce that they have shifted. They just quietly start behaving a little differently. Sometimes better, sometimes worse, sometimes the same words now need a different setting to land. If you never re-test, you can spend months running a workflow that used to be excellent and is now merely fine, without ever noticing the slide.

Thirty minutes a week keeps you learning. This monthly check keeps you from drifting.

How to do it:

1. Once a month, open your personal file.
2. Pick your two or three most important, most-used prompts or workflows.
3. Run each on a fresh real example and judge the result with clear eyes.
4. If something got worse, adjust it and update the file. If a new feature would do the job better, switch and note that too.
5. While you are there, check whether you are still using the right tool and the right model for each job. Defaults change.

A professional example: Tom, a project manager, had a beautifully tuned prompt that turned his messy meeting notes into clean action items with owners and due dates. It ran flawlessly for half a year. During a monthly re-test he noticed the new model version was now occasionally inventing a due date he never mentioned. A two-minute tweak, adding "only include dates I explicitly stated, otherwise mark TBD," fixed it and made the prompt better than before. Had he not re-tested, he would have kept shipping confident action lists with phantom deadlines, and the error would have been his name on it, not the model's.

Put the monthly re-test on your calendar as a recurring event. It is the smallest possible insurance policy against the thing you actually fear, which is silently falling behind without knowing it.

---

5. Teach one thing to a colleague

The move: Once a week or so, teach one small AI thing you have learned to a coworker. A two-minute "here, watch this" at someone's desk. That is the whole step, and it does more for your own mastery than another hour of reading ever could.

There is an old truth in every field. You do not really understand something until you can explain it to someone else. Teaching forces you to organize what you half-know into something clear, and the gaps in your own understanding light up the moment a colleague asks "wait, why does that work?" The act of answering is what turns a loose trick into solid, owned knowledge.

This step also quietly builds something valuable for you professionally. The person who teaches the team becomes, without trying, the person the team sees as the one who gets it. That reputation compounds the same way the skill does.

How to do it:

1. Pick one small thing from your week that worked. One prompt, one feature, one habit.
2. Find a colleague who would benefit and show them, on a real example, in under five minutes.
3. Let them try it themselves while you watch, so it sticks for them too.
4. Notice the questions they ask. The ones that stump you are your next week's experiment.

A professional example: Aisha, an HR business partner, learned to have the AI role-play a tough conversation before she had it for real. The next day she showed a manager who was dreading a performance discussion how to do the same in five minutes. Explaining it forced her to articulate exactly why the role-play helped, which sharpened her own use of it. The manager now swears by it, tells other managers, and credits Aisha. One five-minute share turned a private trick into team-wide practice and made Aisha the obvious person to lead her department's AI rollout.

If you work alone or remote, teaching still works. Post the tip in a team channel, write it up in two sentences, or explain it to a peer in another company. The teaching is the lock. The audience is secondary.

The one-page cheat sheet

Print this. Pin it near your screen. The whole operating system fits on one page.

Every week, about 30 minutes:

1. **Test one new thing on one real task.** Pick real work you have to do anyway, do one part of it a new way, keep what beats your old method. Fifty-two small experiments a year is how fluency is actually built.
2. **Save what works to your personal file.** One document. Every time a prompt or workflow lands well, paste it in with a one-line note. Glance at it before you start any task. It is the one asset that follows you across every tool and model change.
3. **Skim your two or three trusted sources.** Vendor changelogs for truth (Claude, ChatGPT, Gemini, or Copilot release notes, whichever you use). One human filter for meaning (Ethan Mollick's "One Useful Thing"). Optionally one daily aggregator (TLDR AI or The Rundown, pick one). Ignore the rest of the firehose.
4. **Teach one thing to a colleague.** Two minutes, a real example, let them try it. Explaining it is what turns your trick into mastery and turns you into the person who gets it.

Once a month, a bit longer:

5. **Re-test your top tools and prompts.** Re-run your two or three most important prompts on fresh examples. Models change silently. Catch the drift, update your file, confirm you are still using the right tool and model for each job.

The hype filter, any time you see a big claim:

- Who is saying it, and what are they selling?
- Is there a primary source, or is it a rumor?
- Does it match what I see when I actually try it?
- Is this genuinely new, or recycled for clicks?

That is the entire system. It is deliberately small, because a routine you will actually run beats an ambitious one you abandon in a month. The fear of falling behind dissolves the moment you have a reliable way to stay current. You just installed one.

This is a companion to The Artificial Advantage: Mastering AI in the Modern Workplace by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%. This guide keeps that edge sharp for good. Get the book at elvinpeters.com.

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google, OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.