

TELL ME THE TRUTH

The Anti-Hallucination Field Guide: How to Stop AI From
Making Things Up

The truth-first system professionals use to catch AI when it lies, force it to
cite its sources, and trust the answer before they put their name on it.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

There is one fear that stops smart professionals from leaning on AI, and it is the right fear to have. You cannot trust it. You ask a clean question, you get back a confident, polished, well-written answer, and you have no idea whether half of it is invented. So you either waste the time you saved double-checking everything, or you paste it into a client email and pray.

That fear is well founded. AI models do make things up. The industry word for it is "hallucination," and it is not a rare glitch. It is a predictable feature of how these tools work. Here is the part nobody tells the non-technical user: it is also largely controllable. The same companies that build these models publish exact instructions for forcing accuracy out of them, and almost none of those instructions have reached the people who need them most.

This guide is the part of the manual that deals with trust. It explains, in plain English, why the machine invents things, the handful of tasks where it lies the most, and the word-for-word moves you can make to force it to tell you the truth or admit that it does not know. None of this requires code. All of it is drawn from the published guidance of Anthropic, OpenAI, and Google, the companies that actually build these systems. By the end you will stop hoping the answer is right and start being able to check.

1. Why it makes things up, in plain terms

To catch a liar you have to understand why it lies, and the AI is not lying the way a person does. It has no intent to deceive. It was simply built to do something other than tell the truth.

A language model is a very sophisticated prediction engine. Its entire job is to look at the words so far and produce the most plausible next words. Plausible and true are not the same thing. When the model knows the answer, the most plausible continuation happens to be correct. When it does not know, the most plausible-sounding continuation is still produced with the exact same confidence, because the machine has no internal signal that says "I am now guessing."

In September 2025, researchers at OpenAI published a paper explaining a second, deeper cause. Models are trained and graded the way students are graded on a multiple-choice exam. On that kind of test, a blank answer scores zero and a wild guess might get lucky. So the smart test-taking strategy is to always guess and never leave anything blank. The researchers found that AI models have been rewarded, over and over during training, for

confident guessing rather than for honestly saying "I don't know." We trained the bluffing in. As their paper put it, the systems are optimized to be good test-takers, and guessing when uncertain improves the test score.

Two things follow from this, and both are useful.

First, a wrong answer and a right answer look identical on the surface. The confidence in the writing tells you nothing about the accuracy of the content. Stop reading fluency as a signal of truth.

Second, because the bluffing is baked into the default behavior, you can suppress it by changing the instructions. If you explicitly tell the model that admitting uncertainty is acceptable, you remove the pressure to guess. That single move, which the next section covers, is the highest-leverage thing in this entire guide.

What the machine is really doing: A model that says "The contract auto-renews on a 90-day notice period" when the contract says no such thing is not lying to you. It generated the most contract-sounding sentence it could, and no part of its training stopped it. Your job is to put that stop in yourself.

2. Give it permission to say "I don't know"

This is the simplest fix and the one almost nobody uses. By default the model feels pressure to produce an answer. Take that pressure off and the false answers drop sharply.

Anthropic's own guidance on reducing hallucinations lists this first. The instruction is to explicitly give the model permission to admit uncertainty, and they note that this simple technique can drastically reduce false information. You are not asking it to be smarter. You are telling it that "I don't have enough information" is a winning answer instead of a failing one.

How to do it: Add one sentence to any factual request. Tell it that if it is unsure or the information is not available, it should say so plainly rather than guess.

Before: "What were the key terms of the 2023 supplier agreement with Northwind Logistics?"

After: "What were the key terms of the 2023 supplier agreement with Northwind Logistics? If you do not have reliable information on this specific agreement, tell me that directly instead of guessing. Do not invent terms, dates, or figures."

The first version invites a fabricated, plausible-looking contract summary. The second version gives the model an honorable exit, and a well-built model will take it. You would rather hear "I can't confirm that" in three seconds than discover it in front of a client.

3. Make it quote the source before it answers

This is the single most powerful technique in the guide, and it is the one professionals working with documents rely on most. The principle is "quote first, answer second." Before the model is allowed to give you its conclusion, you make it pull the exact words from the source that it is relying on.

Anthropic documents this directly. For tasks involving long documents, their guidance is to have the model extract word-for-word quotes first, before performing the task, which grounds its responses in the actual text and reduces hallucinations. The reason this works is mechanical. If the model has to find and copy a real sentence from your document, it cannot answer from a vague impression. It is forced to point at something that actually exists. When there is no supporting sentence to point at, the gap becomes visible to both of you.

How to do it: When you give it a document to work from, split the task in two. Step one is to quote. Step two is to answer using only those quotes.

Before: "Based on this 30-page lease, can we sublet the office?"

After: "Using only the lease below, answer whether we are permitted to sublet the office. First, copy out the exact clauses that deal with subletting, word for word, with their section numbers. Then, using only those quoted clauses, give me a plain-English answer. If there is no clause that addresses subletting, say 'The lease does not appear to address this' rather than inferring."

The difference is night and day. The "before" gets you a confident yes or no you cannot verify. The "after" gets you the actual clauses, in the lease's own words, that you can read with your own eyes before you trust the conclusion. You are no longer trusting the machine. You are checking its work.

4. Make every claim auditable with a self-check

The previous move grounds the answer before it is written. This one cleans up the answer after it is written. You ask the model to go back through its own response, find support for each claim, and delete anything it cannot back up.

Anthropic's guidance describes exactly this. Make the response auditable by having the model cite a quote or source for each claim, then have it verify each claim by finding a supporting quote after it generates the response. If it cannot find a quote, it must retract the claim. You are turning the model into its own fact-checker, and because retracting is now the instruction, it stops defending its guesses.

How to do it: After it produces a draft, run a second pass with a clear rule for unsupported claims.

The follow-up prompt: "Now review the summary you just wrote. For every factual claim, every number, every date, and every name, find the exact supporting line in the source documents and put it in brackets right after the claim. If you cannot find direct support for a claim, delete that claim entirely and mark the spot with empty brackets [] so I can see where it was removed."

What you get back is a draft where every fact carries its receipt, and the holes are marked instead of hidden. The empty brackets are gold. They show you precisely where the model was filling in with invention, which is exactly where your attention belongs.

5. Know the five places it lies most

Hallucination is not evenly spread. The model is highly reliable on some tasks and a confident fabricator on others. Knowing the danger zones tells you when to relax and when to verify every word. These are the five tasks where the bluffing instinct is strongest.

Specific facts about narrow or obscure subjects. The more specific and less famous the subject, the thinner the model's real knowledge and the more it fills the gap with plausible invention. A small company's founding year, a niche regulation's exact wording, a particular person's job history. The OpenAI researchers gave a memorable example: asked for a specific person's dissertation title or birthday, a model will often produce a confident, specific, and completely wrong answer rather than admit it does not know.

Numbers, statistics, and calculations. Models generate text, they do not natively do math, and they will produce a clean-looking figure that is simply wrong. Treat every statistic, percentage, and total as unverified until you have checked it or had the tool show its arithmetic step by step.

Direct quotes and attributions. Asked who said something, or for the exact wording of a quote, models frequently invent the quote, misattribute a real one, or blend two together. This is one of the most common and most embarrassing failures. Never publish a quote the AI gave you without finding the original.

Recent events. Every model has a training cutoff date, a point after which it has learned nothing. Ask about anything more recent and it is working blind. It may tell you confidently about a deal, a launch, or an appointment, and have the details wrong or invented, simply because the real event happened after its knowledge ends. This is the danger zone where a search tool, covered in section seven, becomes essential.

Legal, medical, and financial specifics. This is where a hallucination stops being an annoyance and becomes a real-world risk. Specific case citations, statute numbers, drug interactions, dosages, tax thresholds. The model can produce authoritative-sounding specifics that do not exist. Lawyers have been formally sanctioned for filing briefs that cited court cases an AI invented. In these domains, treat the AI as a drafting assistant whose every factual claim must be confirmed against a real source, never as the authority itself.

The rule of thumb: The more a wrong answer would cost you, and the more specific the fact, the harder you verify. A brainstorm needs no checking. A statute citation in a contract needs every digit confirmed.

6. A verification workflow any non-coder can run

You do not need software or technical skill to fact-check AI output. You need a short, repeatable routine. Here is the one to run before you trust anything that carries consequences. It takes a couple of minutes and it is the difference between using AI safely and getting burned.

Step one. Ask it to separate fact from inference. Before checking anything, find out what actually needs checking.

"Go back through your answer and split it into two lists. List A: statements you are confident are established fact. List B: statements that are your interpretation, estimate, or assumption. Be honest about which is which."

This alone surfaces the soft spots the model was papering over.

Step two. Make it rate its own confidence. Models can flag their own uncertainty when you ask, and it is more reliable than the confident default tone suggests.

"For each fact in List A, rate your confidence as High, Medium, or Low, and tell me which ones I should independently verify before relying on them."

Step three. Verify the Low and Medium items at the source. This is the human step, and it is non-negotiable for anything important. Take the specific names, numbers, quotes, and dates the model flagged, and confirm them against a primary source. The original document. The company's own website. The actual statute or study. Do not confirm one AI by asking another AI. Go to the real thing.

Step four. Cross-examine with a second, independent question. Ask the same factual question again in a completely different way, in a fresh chat with no prior context. If you get two answers that genuinely conflict, at least one is invented, and you have caught it. Consistency is not proof of truth, but a contradiction is proof of a problem.

Before: You read the AI's polished summary, it sounds right, you ship it.

After: You run the four steps, the model admits two of its five "facts" are Medium confidence, you check them, one is wrong, you fix it before anyone sees it. Two minutes spent, one disaster avoided.

7. Use grounded tools, not raw chat, for anything factual

Here is a distinction that changes everything, and most people do not know it exists. There are two very different modes you can run an AI in.

The first is raw chat. The model answers from memory, meaning the frozen snapshot of text it was trained on. This is where hallucination lives, because the model is reconstructing facts from a fuzzy internal impression with no source in front of it.

The second is grounded mode. The tool first goes and retrieves real, current information, a live web search or a specific document you provided, and then answers using that retrieved material. Now the model is reading and summarizing a real source instead of recalling one from memory. The difference in reliability is enormous.

Google describes this plainly in its own product documentation. Grounding a model's responses in real-world information reduces hallucinations, gives access to information

beyond the model's training cutoff, and returns inline citations that link specific parts of the answer back to the actual web sources, so you can see exactly which claim came from which page. That last part is the whole game. A grounded answer comes with its receipts attached.

How to do it, in practice, with no technical setup:

- For anything recent or anything you need to be current, use a tool or mode that searches the web and shows you links. Most major assistants now offer this, sometimes labeled as search, browsing, or research mode. If the answer comes back with clickable source links, you are grounded. If it comes back as bare confident prose with no sources, you are in raw chat and you should not trust the facts.
- For questions about your own material, give the model the actual document rather than relying on its memory of something similar. An answer drawn from the file you pasted is far safer than an answer drawn from the model's general training.
- Always click the sources. A grounded tool can still misread a page, or cite a weak source. The citations exist so you can check them. Open one or two and confirm the answer actually matches what the source says.

Before: "What's the current corporate tax filing deadline for a Canadian small business?" asked in raw chat, answered confidently, and possibly out of date by a year.

After: The same question asked in a search-grounded mode, which returns the answer with a link to the government's own page, which you click and confirm. One is a guess wearing a suit. The other is a sourced fact you can stand behind.

The one-page cheat sheet

Keep this beside you. Stack these moves, do not pick just one.

1. **Fluency is not truth.** A confident, polished answer is not a correct answer. The model writes wrong answers and right answers in the exact same voice.
2. **Give it an exit.** Add "If you're not sure, say so and don't guess" to every factual request. This one line cuts fabrication sharply.
3. **Quote first, answer second.** With any document: "Copy the exact relevant lines word for word, then answer using only those." Forces it to point at real text.
4. **Make it self-check.** After a draft: "Find a supporting quote for every claim. Delete any claim you can't support and mark the spot with []". The empty brackets show you

where it was inventing.

5. **Know the five danger zones.** Obscure facts, numbers, quotes, recent events, and legal/medical/financial specifics. Verify these every time.
6. **Run the four-step check.** Separate fact from inference, have it rate its own confidence, verify the shaky items at the original source, then re-ask in a fresh chat and watch for contradictions.
7. **Prefer grounded mode.** For anything factual or recent, use a tool that searches and shows source links. No links means raw memory, which means do not trust the facts.
8. **Go to the real source.** Never fact-check one AI with another AI. Confirm against the primary document, the official page, the actual study.

The goal was never to make the AI perfect. The goal is to never be the person who got fooled. These moves put a check between the machine's confidence and your reputation, and that check is the entire difference between using AI like an amateur and using it like a professional.

This is a companion to [The Artificial Advantage: Mastering AI in the Modern Workplace](#) by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%, including the diagnostics that catch a weak answer before it costs you. Get it at elvinpeters.com.

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google, OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.