



Direct Your Tool

Chapter 8 of The Artificial Advantage

A free chapter from the book. Directing AI well is the whole game, and this is where you learn it.

ELVIN PETERS

"I suppose it is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail."

— Abraham Maslow, *The Psychology of Science*

Marcus's first serious AI session lasted forty-five minutes and produced almost nothing he could use.

He was building a pitch deck for a product launch. Real stakes, senior leadership in the room. He'd read somewhere that the reasoning model was the most powerful option, so he opened it and handed it the brief: market context, competitive positioning, and the three scenarios he wanted leadership to choose between.

The model went to work, slowly. It returned structured analysis of each scenario, cross-referenced assumptions, uncertainty flags throughout. Careful. Thorough. And completely wrong for what he needed.

He needed slides. The reasoning model had treated his brief as an analytical puzzle and handed back a fifty-page strategy document.

A colleague two desks over watched him scroll, and scroll, and scroll. She asked what he was building. A deck, he told her, for Thursday. She pointed at his screen. "You used the reasoning model for a creative task." Then she showed him the fix: switch to the default model, and add one line to the end of his brief, format this for a ten-slide executive presentation, one key idea per slide. Six minutes later he had a working skeleton.

The model wasn't the problem; the match was. Marcus had reached for raw power when the task called for a different mode, and that single mismatch cost him most of an hour. By the end of this chapter you'll never make it.

The Dropdown Most Beginners Ignore

Every AI tool shows a dropdown listing several model options. Most beginners do one of two things with it: ignore it completely, or burn twenty minutes researching which model is "best" before typing a single word. Neither is the move. Those options are versions tuned for different kinds of work, and the right pick depends entirely on what the task in front of you actually requires.

It helps to think in terms of three kinds of thinking you can ask for, rather than three separate buttons. There's fast thinking (quick, broad, good enough for most things), deep thinking (slow, sequential, shows its work), and divergent thinking (deliberately unexpected). You summon them through two mechanisms. One is the model dropdown,

where a "default" model gives you fast thinking and a "reasoning" model gives you deep thinking. The other is your prompt, which can pull deep or divergent thinking out of the default model with a single line.

That overlap raises a fair question, so here's the rule that resolves it. Switch models when the task is genuinely long or when the methodology has to be auditable, the kind of thing someone might challenge line by line. Prompt for depth instead when you're on a free tier, when speed matters, or when the task is short. Most days, the prompt is enough. The dropdown is for the heavy, high-scrutiny work.

One note before the specifics. The model names below are examples, current as I write this, and they will change every release. What persists is the three kinds of thinking. Learn the distinction and the labels will sort themselves out.

The Default Model

This is where most of your work belongs. The default model balances intelligence and speed for the broad run of tasks: emails, summaries, brainstorming, quick research, explanations, first drafts, planning outlines.

How much of your work? Call it a rule of thumb, the 80/20 of AI use: the large majority of what you do never needs anything heavier than the default. Treat that as a planning heuristic, the same shape as the Pareto idea, rather than a measured statistic. The exact percentage is beside the point. What matters is that heavier rarely means better, and most tasks that feel like they demand the reasoning model just need a sharper prompt to the default one.

The Reasoning Model

Every platform has its own version and its own name for it: OpenAI's o-series, Anthropic's extended-thinking toggle, and others, with the labels shifting every release. The idea has held steady. It's a model that reasons step by step before it answers, breaking a problem into parts and checking its own assumptions instead of firing off the first thing that comes to mind. The answer arrives later because the model is genuinely working through the problem.

Reserve this mode for tasks where getting the sequence of reasoning right matters as much as the final answer. Picking which of three vendors to go with and being able to defend the ranking when your boss pushes back. Planning a project where step three depends on how step one resolves. Financial or legal analysis. Debugging a stubborn

piece of code. The common thread is that the logic has to hold up to scrutiny, not just sound right.

Two limits are worth knowing. Reasoning mode burns more tokens, which eats into your usage on free and lower-paid tiers. And on a task that doesn't call for deliberate sequential reasoning, all that extra processing buys you nothing over a well-prompted default model. Using it to draft a two-line email is the hammer meeting a task that was never a nail.

Creative Mode

This kind of thinking doesn't live in a dropdown. It lives in the prompt. The default model reaches for the statistically likely response. Creative mode reaches the other way, toward the contrarian take, the unconventional angle, the idea a standard request would never surface. You turn it on by asking for it, in plain words:

"Give me the unexpected version."

"Ignore the obvious answer."

"What would a contrarian say?"

"Ten ideas, where the last three are genuinely strange."

Here's the gap it opens, using Marcus's deck. He needed an opening line for the title slide and asked the default model for "three taglines for our launch." Back came the expected: "Introducing [Product]: The Future of Productivity," "Work Smarter, Not Harder," "The Tool Your Team Has Been Waiting For." Serviceable, forgettable, the kind of line every deck in the building already uses. Then he reframed it: "Give me three opening lines a competitor would never dare to put on a slide, including one that's a little uncomfortable." The third option made the room laugh and then go quiet: "Most of your team is already doing this badly in secret. Let's make it official." That became the hook. Same model, same minute. The first request got him the average of every launch deck ever written. The second got him the one nobody else would have said. Most professionals never see this gear, because they take the first response and move on. The ceiling is much higher than the default shows you, but you have to ask for it.

Reasoning Depth Without Switching Models

Here's the part that matters most if you're on a free tier, or just don't want to bother switching models. You can pull reasoning-mode depth out of the default model with a single line. Add any of these to a standard prompt:

"Think step by step."

"Show your reasoning before giving your answer."

"List your assumptions first."

"Walk me through how you arrived at this."

Researchers call this chain-of-thought prompting, and it works for a simple reason. Asking the model to reason out loud before concluding produces more careful, more transparent output, the same thing you'd ask of a sharp colleague walking you through their thinking. The same model, handed the same task, gives you a fundamentally different answer depending on whether you ask for a quick list or ask it to show its reasoning first.

Marcus: The Next Morning

The lesson from the deck session was simpler than "the reasoning model is useless." Marcus had pointed it at the wrong task.

His launch deck included a competitive analysis section that genuinely needed careful, multi-step logic, methodology that would hold up under executive questioning. So the next morning he switched the reasoning model back on, but only for that section. He fed it the competitor data, the twelve evaluation dimensions, and the weighting rationale, and asked it to work through the scoring step by step before summarizing. Flag any dimension, he added, where the evidence runs thin.

Forty seconds. Two flags he'd missed in his own review. He fixed both before the deck left his desk. The rest of it ran on the default model: six minutes for the skeleton, twenty minutes of editing.

That was the real shift. Model selection was never a setting he chose once at the top of a session. It was a call he made section by section, sometimes paragraph by paragraph. Jordan, from the previous chapter, worked the same way inside a single client pitch: default model for the narrative framing, reasoning model for the scoring matrix where the methodology had to survive a hard look. The choice happens at the task level, down where the work actually lives.

Your Leverage Stack

Stack a chosen tool on top of a directed mode and you've assembled exactly what the leverage argument earlier in the book was pointing at: a professional aiming the right instrument, in the right configuration, at the task in front of them.

The multiplier was never having AI. It comes from knowing how to direct it. The professional who keeps creative work on the default model, switches to reasoning when the methodology has to hold, and prompts for the contrarian take when the first answer isn't good enough operates a full level above someone who accepts whatever the default hands them. The gap looks intimidating from the outside. It comes down to three habits, and Marcus closed it in a single morning. Underneath all three is the same quiet act: deciding what the task in front of you actually needs before you reach for raw power.

Exercise: One Decision, Two Modes

Pick a real decision you're weighing right now: a professional tradeoff, a planning question, a judgment call you've been putting off.

Send this to your AI tool: "What should I consider when [your situation]?" Read the response.

Then send the same prompt with this added: "Think through this carefully before answering. Consider the tradeoffs, the second-order effects, and what I might be missing."

Compare the two. The first shows you the default model's reach. The second shows you reasoning-mode depth pulled from the very same model, with nothing but a line of text. Whether the extra depth was worth the extra time is your call, but now you're deciding with the guesswork stripped out.

Key Takeaways

Three kinds of thinking live inside every AI tool: fast (the default model), deep (the reasoning model), and divergent (creative prompting). You summon them through the model dropdown or through your prompt.

The default model handles the large majority of professional work. Treat "most of it" as a planning rule of thumb, not a measured number, and reach for the reasoning model only when the logic has to survive scrutiny.

Creative mode ships with every model but rarely gets switched on. Ask for the unexpected version, the contrarian take, the idea past the obvious, and the ceiling rises well above the default answer.

Chain-of-thought prompting, adding "think step by step" or "show your reasoning first," pulls deep thinking out of the default model. No switch required, and it works on free tiers.

Mode selection is an ongoing judgment, not a one-time setting. Choose at the task level, section by section, rather than once at the top of a session. Each choice is a small act of knowing what the work in front of you actually needs.