

WHAT NOT TO TYPE

AI Security and Privacy for Professionals: What's Safe
to Share, and What Isn't

The clear line between safe and not safe, so you can use AI on real work
without leaking a client, a contract, or your job.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

There is a quiet question sitting behind almost every serious use of AI at work. You have the contract open, or the client list, or the financials, and you are about to paste it into a chat box, and a small voice says: *wait, is this even allowed?*

Most people resolve that question one of two ways, and both are wrong. Some paste everything and never think about it again, which is how confidential material ends up somewhere it should not be. Others get scared and refuse to use AI on anything that matters, which means they only ever use it for toy tasks and never get the real value. The truth lives in between, and it is more reassuring than either camp expects. There is a clear, knowable line between what is safe to put into these tools and what is not. Once you can see that line, the fear goes away and the caution stays, which is exactly the combination you want.

This guide draws on the actual published data-handling policies of the four companies most professionals use: Anthropic, which makes Claude; OpenAI, which makes ChatGPT; Google, which makes Gemini; and Microsoft, which makes Copilot. The policies are public. Almost nobody reads them, because they are long and written for lawyers. What follows is the plain-English version, current as of 2026, with a realistic professional example for every rule so you can see exactly how it applies to your day.

One promise up front. This is not a lecture about how AI is dangerous. The tools are genuinely useful on sensitive work when you use the right version and a few simple habits. The goal here is to make you the person in the room who actually knows the answer when someone asks "are we allowed to do this?"

1. Know which version you are using, because it changes everything

This is the single most important thing in this guide, and it is the thing people understand least. The same chat box can have completely different rules underneath it depending on which account it is attached to. There are, broadly, two worlds.

The first is the **consumer world**: the free version and the cheap personal subscription you signed up for with your own email. The free Gemini app, a personal ChatGPT login, the Claude app on a personal plan, the consumer Copilot. The second is the **business and enterprise world**: a paid work account your employer set up, an enterprise or team plan, or the developer API that sits behind a company's own tools.

The difference matters because of one thing: whether your conversations can be used to train the AI, and who might see them. In the business and enterprise world, the answer across all four major vendors is consistent and clear. By default, OpenAI does not train on inputs or outputs from ChatGPT Team, ChatGPT Enterprise, or the API. Anthropic's commercial terms cover Claude for Work, Team, Enterprise, and the API, and that data is not used for model training and is not subject to any opt-in toggle. Google commits that Workspace business data is not used to train its core models and is not subject to human review. Microsoft states that prompts, responses, and the data Copilot reaches through your organization's files are not used to train its foundation models. Business-grade accounts are private by design.

The consumer world is a different story, and it has shifted recently. On a personal ChatGPT account, your chats can be used to train future models unless you go into settings and turn that off. Anthropic changed its consumer terms in 2025 so that personal Claude users now have to make an active choice about whether their conversations help train the model, and if you allow it, the retention period extends to five years. The free Gemini app may use your conversations to improve the service, and a slice of conversations gets reviewed by humans for quality, with reviewed chats kept for up to three years even though Google tries to disconnect them from your account first.

The professional takeaway: before you paste anything that matters, know which world you are in. If you are doing real work with real client or company information, you want to be in the business or enterprise world, on an account your employer controls. If all you have is a personal login, treat it like a public space.

A realistic example. A management consultant has both a personal ChatGPT subscription she pays for herself and a company ChatGPT Enterprise seat. She drafts a generic framework for a blog post in her personal account, which is fine because there is nothing confidential in it. When she moves to analyzing a client's actual restructuring plan, she switches to the company enterprise account, because that is the one with the contractual guarantee that the data will not be used for training and will not leak into the next version of the model. Same tool, same screen, two completely different sets of rules. She knows which is which.

2. Learn what actually counts as sensitive

People tend to picture "sensitive information" as something dramatic, like a password or a leaked memo. In practice the sensitive material in your day is more ordinary and more

common than that, which is exactly why it slips through. Here is the working list.

Personal information about other people. Names tied to anything private. A client's home address, a customer's phone number, an employee's salary, a patient's diagnosis, a job applicant's details. This is the category regulators care about most, and it is the one professionals leak most often, usually by pasting a spreadsheet or an email thread without thinking about the names inside it.

Financial details that are not yours to share. A client's revenue figures, unannounced numbers, deal terms, account numbers, anything that would move a stock or break a deal if it got out.

Trade secrets and proprietary work. Source code, unreleased product plans, pricing models, the secret recipe, the strategy deck that has not been presented yet. The things that are valuable precisely because competitors do not have them.

Credentials. Passwords, API keys, access tokens, the contents of a `.env` file. These should never go into a chat box, ever, on any tier, full stop. A credential pasted into a prompt is a credential you now have to assume is compromised.

Regulated categories. Health information, legal matters under privilege, anything covered by a specific law or contract. In Canada this means PIPEDA and provincial health privacy rules. If you are dealing with US clients it may mean HIPAA. In Europe it means GDPR. You do not need to memorize the acronyms. You need to recognize that health, legal, and certain personal data come with extra rules attached.

A realistic example. A bookkeeper wants help cleaning up a messy client invoice spreadsheet. The spreadsheet seems harmless, just numbers and dates. But three columns in, it has the client's customers listed by full name, with their outstanding balances and, in the notes field, two lines about why one customer is disputing a charge. That is personal information and financial information about third parties who never agreed to have their data sent anywhere. The spreadsheet looked like "just numbers." Most sensitive data is hiding in something that looked like just numbers.

3. Anonymize before you paste

Here is the most freeing idea in this whole guide. The AI almost never needs the real names to do the job. It needs the structure, the numbers, the relationships, the problem. The identities are usually decoration that you can strip out without losing any of the help.

So strip them out. Replace the real client name with "Client A." Replace the real person with "the customer." Swap exact figures for rounded or relative ones if the exact figures are what is sensitive, or keep the figures and remove the names if the names are the sensitive part. You are not deceiving the tool. You are giving it everything it needs to be useful while keeping the part that would cause harm safely on your own machine.

This single habit lets you use even a consumer-tier tool on work that touches sensitive material, because by the time the words reach the chat box there is nothing identifying left in them. It is the difference between sending a redacted document and sending the original.

A realistic example. A lawyer wants help drafting a firm but polite letter to the other side in a contract dispute. She does not paste the real letter with both companies named, the dollar amount, and the privileged details. She writes the prompt as: "I represent a mid-sized manufacturer in a dispute with a supplier. The supplier delivered late and we are seeking a credit. Help me draft a firm, professional demand letter. Keep names and figures as placeholders I will fill in." She gets a polished draft in two minutes, and nothing privileged ever left her control. She fills in the real names in her own word processor afterward.

4. Summarize and quote, don't paste the whole thing

There is a difference between pasting an entire confidential document into a tool and asking the tool to help you think about it. Often you can get what you need by bringing the AI a summary instead of the raw material.

Rather than pasting the full 40-page contract, you can paste the three clauses you are actually worried about, with the names removed. Rather than uploading the entire HR file, you can describe the situation in general terms and ask for an approach. Rather than dropping in the whole customer database, you can paste ten anonymized example rows and ask it to write the formula, then run that formula yourself on the real data on your own computer.

This does two things at once. It shrinks the amount of sensitive material that leaves your machine down to the minimum the task requires, and it usually produces a sharper answer, because you have done the work of pointing the AI at the part that matters instead of making it hunt through pages of irrelevant text.

A realistic example. An HR manager is dealing with a difficult termination and wants help wording it carefully. Instead of pasting the employee's full file, the performance reviews, and the medical note that explains the background, she writes: "I need to word a termination conversation for an employee whose performance has been documented as below standard over six months, after a formal improvement plan. Help me script it to be clear, humane, and legally careful. Assume Ontario employment law." She has the situation, none of the identity, and none of the medical detail that was never relevant to the wording anyway. The answer is just as good and the file never left her office.

5. Match the tier to the task

You do not need to use the most locked-down tool for everything. That would be exhausting and most of your work does not require it. The skill is matching the sensitivity of the task to the privacy of the tool, and there are really only three levels.

Public or generic work can go anywhere, including free consumer tools. Brainstorming blog topics, drafting a generic template, rewriting a paragraph for tone, explaining a concept to yourself. There is nothing here that could harm anyone if it were seen, so there is nothing to protect.

Internal but not catastrophic work belongs on a business-grade account that your employer controls, where the data is not used for training. Most day-to-day professional work lives here. Drafting client emails, analyzing your own company's internal reports, working through a plan.

Highly sensitive or regulated work belongs on an enterprise account with the strongest available settings, and sometimes on tools your IT department has specifically approved, possibly with zero-retention agreements in place. Anthropic and OpenAI both offer zero-data-retention arrangements for enterprise API customers, where the inputs and outputs are not stored beyond what is needed to screen for abuse. This is the tier for the merger, the lawsuit, the health records, the crown-jewel source code.

A realistic example. A financial advisor uses the free tier of a chatbot to draft a general explainer about how RRSPs work, because that information is public and helps nobody to keep secret. He uses his firm's enterprise account to analyze an anonymized version of a client's portfolio allocation, because that is real work on real numbers and he wants the training-excluded guarantee. And when the compliance team needs to review actual client records against a regulation, that runs only on the specific platform the firm's IT

department vetted and approved. Three tiers, three tools, one clear logic connecting them.

6. Read the room: company policy and compliance

Even when a tool is technically safe, your employer or your client may have rules about it, and those rules win. A surprising number of AI privacy incidents are not really about the AI at all. They are about someone using a tool their company never approved, or feeding it data a contract said they would protect.

So before you build AI into how you work, find out three things. First, does your company have an AI policy, and what does it say you can and cannot put into these tools? Many do now, and many people have never read theirs. Second, are there approved tools, and are you using one of them rather than a personal account you signed up for on the side? Third, do any of your client contracts include confidentiality or data-handling clauses that would cover the material you are about to paste? A standard nondisclosure agreement often does, even if it never mentions AI by name.

If your company has no policy yet, that is not permission, it is a gap. The safe default in a gap is to behave as though sensitive data should not leave approved, business-grade tools, and to ask before you assume.

A realistic example. A marketing lead at an agency starts using a slick new AI tool she found online to speed up client campaign work. It is a good tool. The problem is that her agency's contract with the client says client materials may only be processed on systems the agency has security-reviewed, and nobody reviewed this one. The tool was never the risk. The unreviewed tool handling contractually protected material was the risk. A two-minute check with her operations manager would have caught it, and now it is the kind of thing that surfaces in an audit. She starts every new tool with that two-minute check from then on.

Safe vs not safe, at a glance

Generally safe to put into any AI tool, including free consumer versions:

- Public information you could find on a website

- Generic templates, frameworks, and how-to questions
- Your own writing that contains nothing confidential, for editing or rephrasing
- Hypothetical or anonymized scenarios with all real identities removed
- Brainstorming, learning, and explaining concepts to yourself

Safe only on a business or enterprise account your employer controls:

- Internal company documents, reports, and plans
- Client work where you have removed or anonymized third-party identities
- Draft communications referencing real but non-public business details
- Analysis of your own organization's non-public numbers

Not safe to paste into a general AI tool at all, regardless of tier:

- Passwords, API keys, access tokens, or anything from a credentials file
- Other people's personal information you have no clear right to share onward
- Privileged legal material or protected health information, unless on a specifically approved and compliant platform
- Trade secrets and unreleased crown-jewel material, unless your IT department has approved the exact tool and settings
- Anything a client contract or company policy forbids, even if the tool itself is secure

The one-page cheat sheet

Pin this next to your screen. Run it as a quick mental pass before you paste.

Before you paste, ask:

1. **Which account am I in?** Personal consumer login, or a business account my employer controls? If it is real work that matters, it belongs on the business account.
2. **Is there anything sensitive in here?** Names tied to private details, financials that are not mine to share, credentials, health or legal material, trade secrets. Look hard, because it usually hides inside something that looked like "just numbers."
3. **Can I anonymize it?** Swap real names for "Client A" and "the customer." The tool almost never needs the real identities to help.
4. **Can I summarize instead of paste?** Bring it the three clauses, not the forty pages. Bring it ten example rows, not the whole database.

5. **Does the tier match the task?** Public work anywhere, internal work on a controlled business account, highly sensitive work only on an approved enterprise setup.

6. **Am I allowed?** Company policy, approved-tool list, client confidentiality clauses. If there is no policy, treat that as a reason to ask, not as permission.

Two rules that never bend:

- Never paste a password, key, or token into any chat box, on any tier, for any reason.
- When you genuinely do not know whether something is allowed, anonymize it or ask first. The minute it costs you is nothing against the cost of getting it wrong.

The professionals who use AI on serious work without ever causing an incident are not the cautious ones who avoid it, and they are not the reckless ones who paste everything. They are the ones who learned this line once and then stopped worrying, because they could see exactly where it was. You can be that person now.

This is a companion to [The Artificial Advantage: Mastering AI in the Modern Workplace](#) by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%, safely and with confidence. Get it at elvinpeters.com.

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google, OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.