

BUILD YOUR OWN AI ASSISTANT (NO CODE)

Build a Reusable, No-Code AI Assistant for Your Exact
Job

Teach the AI your job once, then reuse it forever, with no code and tools you
already pay for.

ELVIN M. PETERS

A companion to The Artificial Advantage · elvinpeters.com

Why this exists

Here is the quiet tax almost everyone pays. You open a fresh chat, and before you can ask for anything useful, you have to re-explain yourself. Who your company is. Who the audience is. The tone your boss likes. The format the report has to land in. The product you sell and the words you are not allowed to use. You type all of it again, get a decent answer, close the tab, and tomorrow you do the whole thing over.

That is the single biggest hidden waste in how most professionals use AI. They treat a tool with memory like a tool without one. They start from zero every single time.

There is a better way, and it has been sitting in plain sight inside the three tools you most likely already have. Claude calls it a Project. ChatGPT calls it a Custom GPT. Gemini calls it a Gem. The names are different and the buttons live in different places, but they all do the same thing. You set up a reusable assistant once. You hand it standing instructions and your reference files. From then on, every time you open it, it already knows your job. You skip straight to the work.

None of this requires code. None of it requires a developer. It is point, click, type, and save. This guide shows you exactly how to build one, walks you through four real examples you could stand up this afternoon, and shows you how to test it so it actually behaves, and how to hand it to your whole team. Everything here is drawn from the official documentation each company publishes, current as of 2026.

1. What these things actually are, in plain terms

Picture the difference between a temp who shows up not knowing anything about your office, and a colleague who has worked beside you for two years. You can ask the colleague for "the usual report" and they just do it. The temp needs the full briefing every time.

A normal AI chat is the temp. A custom assistant is the colleague.

Mechanically, all three tools wrap two things around the same AI you already use:

Standing instructions. A block of text the assistant reads silently before every single conversation. This is where you bake in who it is, who it speaks to, the rules it follows, and the format it returns. You write it once. It applies forever, to every chat you start inside that assistant.

A knowledge base. A set of files you upload once and the assistant can refer to whenever it answers. Your brand guide. Your pricing sheet. Your policy handbook. Your past reports

it should imitate. You stop pasting these into the chat every time, because they live inside the assistant permanently.

That is the whole idea. Instructions tell it how to behave. Files give it the source material to behave with. The official help docs from all three companies draw exactly this line, and once you see it, the tools stop feeling mysterious.

A quick map of where to find each one, accurate as of 2026:

- **Claude Projects.** In Claude, look for "Projects." Even free accounts can create them (up to five), with custom instructions and an uploaded knowledge base. Paid plans get a far larger knowledge capacity. Sharing with teammates is available on the Team and Enterprise plans.
- **ChatGPT Custom GPTs.** In ChatGPT, the builder lives under "My GPTs" or "Explore GPTs," then "Create." You fill in a Configure tab with Name, Instructions, and Knowledge. You can attach up to 20 files. Building a Custom GPT requires a paid plan (Plus, Team, or Enterprise as of 2026).
- **Gemini Gems.** In Gemini, find "Gems," then create a new Gem with a name, instructions, and optional knowledge files like PDFs, Word docs, or Google Drive files. Gems are available on personal accounts, and Google now lets you share a Gem with others through a link.

Pick whichever you already pay for. The skill transfers between all three because the underlying concept is identical.

2. The core move: teach it once, reuse it forever

Here is the mental shift that makes the rest easy. You are not writing a prompt. You are writing a job description for a permanent member of staff.

When you sit down to set up the instructions, answer five questions in plain sentences and paste the answers in. This is the spine of every good assistant, and the official Gemini guidance frames it almost exactly this way (persona, task, context, format).

1. **Who are you?** "You are a senior internal communications writer for a 400-person logistics company."
2. **Who do you serve?** "You write for busy operations managers who skim on their phones."
3. **What is your standing job?** "You turn my rough notes and updates into clear, scannable announcements."

4. **What are your rules?** "Always lead with the action required. Never use corporate filler. Keep it under 200 words unless I say otherwise."
5. **What should every answer look like?** "Return a subject line, a one-line summary, then the body in short paragraphs."

Then you upload the reference files that this assistant should always have on hand. For the communications writer above, that might be a folder of past announcements that landed well, plus your company style guide.

Save it. You are done. From that moment, every time you open this assistant, you type only the new thing ("here are my notes from the warehouse meeting") and it produces the finished announcement in your house style, for your audience, in your format. The context you used to retype every morning is now permanent furniture.

That is the entire payoff. The setup costs you fifteen minutes once. It saves you a few minutes every day for as long as the job exists.

3. Worked example: the brand voice writer

The pain. Every piece of copy that leaves your company is supposed to sound a certain way. Warm but not cutesy. Confident but never hypey. You know it when you read it, but explaining it to a blank chat box, every time, is exhausting, and the results drift.

The build. Create an assistant and give it this kind of standing instruction:

"You are our brand copywriter. Our voice is warm, plain-spoken, and direct. We write for small business owners who are smart but short on time. We never use hype words like 'revolutionary,' 'game-changing,' or 'unlock.' We never use exclamation marks. We favor short sentences and concrete examples over adjectives. When I give you a draft or a topic, return copy that sounds like the examples in the knowledge base."

Then upload three or four of your best past pieces as the knowledge base, the ones everyone agrees sound exactly right. The assistant now has the rules in the instructions and the proof in the files.

Before. A teammate opens a normal chat and types "write a product announcement for our new scheduling feature." Out comes generic SaaS copy full of "revolutionize your workflow" and three exclamation marks. They spend twenty minutes rewriting it into something that sounds like the company.

After. The same teammate opens the brand voice assistant and types the same sentence. The copy comes back already plain, already warm, already matching the past pieces, with no hype words and no exclamation marks, because those rules and those samples are baked in. They tweak one line and ship it. The voice is consistent no matter who on the team requested it, which is the part a freelancer or a junior hire usually cannot deliver on day one.

4. Worked example: meeting notes into actions

The pain. You finish a call with a page of messy notes. Buried in there are decisions, commitments you made, things other people owe you, and open questions nobody resolved. Untangling that by hand after every meeting is the kind of admin that quietly eats an afternoon a week.

The build. Create an assistant whose entire job is this one transformation. Standing instruction:

"You are my meeting analyst. I will paste raw, messy notes from a meeting. Every time, return four labeled sections in this exact order. DECISIONS MADE. MY ACTION ITEMS, each with an owner of 'me' and a due date if one was mentioned. WAITING ON OTHERS, each with the person's name. OPEN QUESTIONS. If something is ambiguous, list it under OPEN QUESTIONS rather than guessing. Keep each item to one line."

You may not even need a knowledge file for this one. The value is entirely in the fixed instruction and the consistent format.

Before. You reread the notes three times, manually pulling out who owes what, and you still miss the one commitment that surfaces awkwardly two weeks later when someone asks where it stands.

After. You paste the notes into the assistant. It hands back the same clean four-section breakdown every time, so your follow-up email and your task list write themselves. The format never varies, which means after a month you can skim a stack of these and trust that nothing fell through, because the assistant was built to catch exactly those four things and nothing distracts it from the job.

5. Worked example: the policy answer-er

The pain. You are the person everyone asks about the rules. What is our refund window. How much notice does a contractor need to give. What does the travel policy say about booking ahead. The answers all live in documents, but nobody reads documents, so they ask you, and you go digging through a PDF for the tenth time this week.

The build. This is where the knowledge base does the heavy lifting. Create an assistant and upload the actual source documents. The policy handbook. The contract template. The travel and expense guide. Then give it a strict instruction:

"You answer questions using only the documents in your knowledge base. When you answer, first quote the exact line or section you are relying on, then give a plain-English answer. If the documents do not cover the question, say 'That is not covered in the documents I have' and do not guess. Never invent a policy."

That "quote first, and refuse if it is not there" instruction is the difference between a helpful answer-er and a confident liar. It forces the assistant to point at the real text instead of making something up.

Before. A colleague messages you "how long is the parental leave again?" You stop what you are doing, open the handbook, scroll, find it, and type the answer back. Three times a day, every day.

After. You give that colleague the assistant. They ask it directly. It replies with the exact clause quoted, then the plain answer, and when someone asks about something the handbook genuinely does not address, it says so instead of bluffing. You have turned yourself from a human search engine back into a person with an afternoon. On a shareable plan, the whole team queries the same documents and gets the same sourced answer, so the policy stops depending on whether you happen to be at your desk.

6. Worked example: the proposal first-drafter

The pain. Your proposals all follow the same shape. The same sections, the same structure, the same boilerplate about your company and your process. Only the client details and the scope change. Yet you rebuild the scaffolding from scratch each time, or you copy an old one and spend an hour scrubbing out the last client's name.

The build. Create an assistant and upload two or three of your strongest past proposals plus a one-page company overview as the knowledge base. Standing instruction:

"You draft client proposals in our standard structure: Overview, Understanding of Your Needs, Proposed Approach, Timeline, Investment, About Us. Match the tone and structure of the example proposals in your knowledge base. Pull the About Us and process language from those examples. I will give you the client name, their situation, and the scope. Leave any number I have not given you as a clearly marked placeholder like [PRICE TBD]. Never invent pricing or dates."

Before. You open a blank document, hunt down the best old proposal, copy it, and start the careful surgery of replacing every mention of the previous client while making sure you did not miss one.

After. You open the assistant and type the new client's name, their situation, and the scope. Back comes a complete first draft in your exact house structure, with your real About Us language, the previous client scrubbed out entirely, and every missing number flagged as a placeholder so you cannot accidentally send a fabricated price. You spend your hour sharpening the argument instead of rebuilding the skeleton.

7. How to test it before you trust it

A custom assistant is only as good as the testing you do before you lean on it. The fix is simple, and it is the same loop the pros use.

Run it on cases where you already know the right answer. Before you let the policy assistant loose, ask it five questions you know the answer to cold. Did it quote the right section. Did it refuse the one you deliberately asked that is not in the documents. If it bluffed an answer to the trick question, your instruction is not strict enough and you tighten it.

Watch for the two common failures. The first is drift, where it slowly stops following a rule. The second is bluffing, where it answers something it should have refused. Both are fixable in the instructions. If it keeps using a word you banned, the ban is buried too deep in a paragraph, so pull it out and make it its own line. If it keeps guessing, add a blunt sentence like "If you are not certain, say so. Do not guess."

Refine in the instructions, not in every chat. This is the whole point and the part people miss. When the assistant does something wrong, do not just correct it in that one conversation and move on. Go back into the assistant's settings and fix the standing instruction so it never makes that mistake again. Every correction you fold back into the

instructions is a correction you never have to make again. The assistant gets sharper over weeks instead of resetting every morning.

Keep instructions short and ordered. A wall of fifty rules confuses the assistant as much as it confuses a new hire. Lead with the few rules that matter most, keep each one to a line, and cut anything you are not sure earns its place. Tighter instructions behave more reliably than longer ones.

8. Sharing it with your team

A custom assistant that lives only on your machine helps one person. The same assistant shared across a team standardizes how an entire department works, which is where the real leverage is.

The mechanics differ by tool, and here is the honest state of it in 2026:

- **Claude Projects** can be shared with teammates on the Team and Enterprise plans. You share by email, individually or in bulk, and you set whether each person can simply use the project or also edit its instructions and files. That permission split matters, because you usually want everyone using the assistant and only a couple of people allowed to change its rules.
- **ChatGPT Custom GPTs** can be kept private to you, shared by link, or published to your workspace on a Team or Enterprise plan, so a GPT you build can become the standard tool everyone in the department opens.
- **Gemini Gems** can now be shared with others through a link. You create the link in the Gems settings and send it, and the people you choose get access to the same Gem you built.

When you share, think about what you are really standardizing. The brand voice assistant means every piece of copy from every teammate sounds like the company, not like whoever wrote it. The policy answer-er means everyone gets the same sourced answer instead of five slightly different versions depending on who they asked. The proposal drafter means every proposal that leaves the building has the same professional shape. You are not just saving yourself time. You are raising the floor for the whole team to your standard, and you only had to build it once.

One caution worth saying plainly. Anything you upload as a knowledge file is readable by everyone you share the assistant with. Before you share, make sure the files inside it are things that audience is allowed to see. Build a separate assistant for anything sensitive rather than sharing one loaded with documents that should stay restricted.

The one-page cheat sheet

Keep this beside you while you build your first one.

1. **Pick your tool.** Claude Project, ChatGPT Custom GPT, or Gemini Gem. Use whichever you already pay for. The skill transfers.
2. **Write the standing instructions.** Answer five questions in plain sentences: who are you, who do you serve, what is your standing job, what are your rules, what should every answer look like.
3. **Upload the knowledge base.** The reference files it should always have on hand. Past work to imitate, the handbook to quote, the style guide to follow. Upload once, reuse forever.
4. **Force honesty on document tasks.** Add the rule "quote the source first, and if it is not in the documents, say so and do not guess." This is what separates a useful assistant from a confident liar.
5. **Test on known answers.** Ask it five things you already know, plus one it should refuse. Watch for drift and bluffing.
6. **Refine in the settings, not the chat.** When it gets something wrong, fix the standing instruction so it never repeats the mistake. Corrections compound.
7. **Keep instructions short and ordered.** Most important rules first, one line each. Tighter behaves better than longer.
8. **Share to standardize.** Hand it to the team to raise everyone to your standard. Check that every uploaded file is safe for that audience first.

The whole thing rests on one shift. Stop opening a blank chat and re-explaining your job. Teach the assistant once and let it remember. The fifteen minutes you spend building it is the last time you have to explain that context for as long as the job exists.

This is a companion to [The Artificial Advantage: Mastering AI in the Modern Workplace](#) by Elvin M. Peters. The book gives you the full framework for directing AI like the top 10%. Get it at elvinpeters.com.

This guide summarizes publicly available best practices for educational purposes. It is independent and is not affiliated with, authorized by, or endorsed by Anthropic, Google,

OpenAI, or Microsoft. All product names and trademarks are the property of their respective owners.